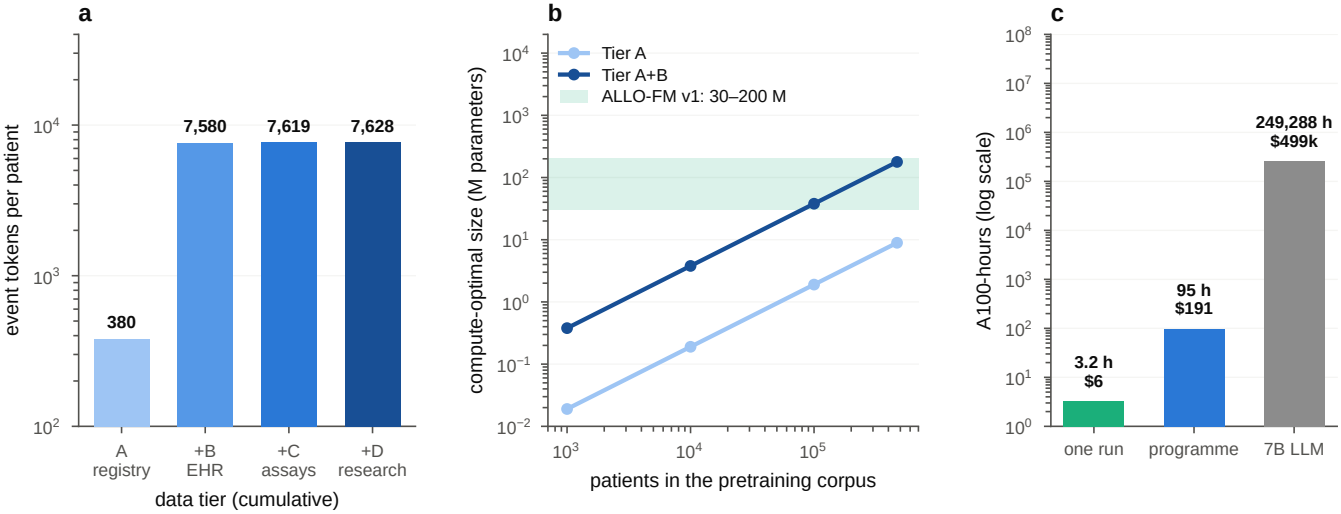
a
event tokens per patient
10^4
10^3
10^2
380
7,580
7,619
7,628
A
registry
+B
EHR
+C
assays
+D
research
data tier (cumulative)
b
compute-optimal size (M parameters)
10^4
10^3
10^2
10^1
10^0
10^{-1}
10^{-2}
Tier A
Tier A+B
ALLO-FM v1: 30–200 M
10^3
10^4
10^5
patients in the pretraining corpus
c
A100-hours (log scale)
10^8
10^7
10^6
10^5
10^4
10^3
10^2
10^1
10^0
3.2 h
$6
95 h
$191
249,288 h
$499k
one run
programme
7B LLM