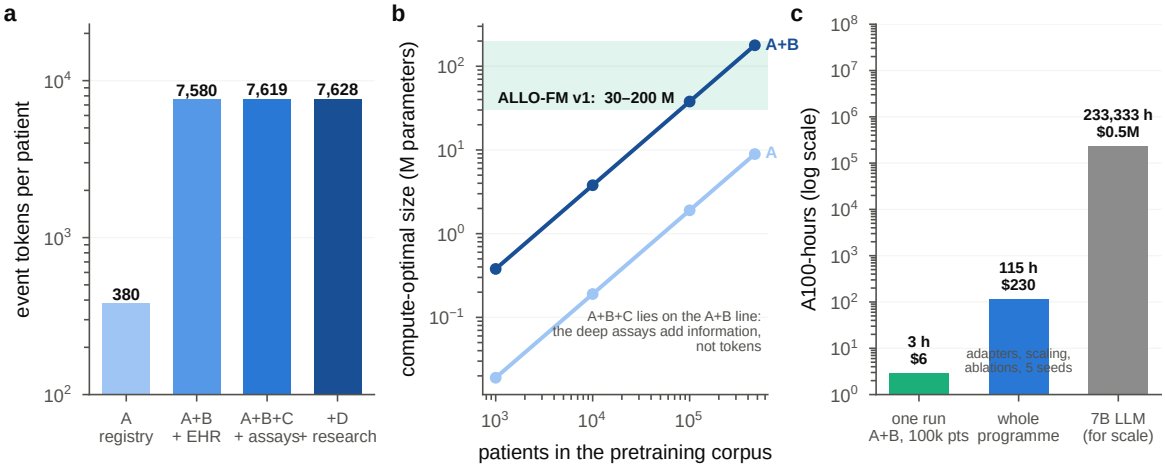
a
event tokens per patient
10^4
10^3
10^2
380
7,580
7,619
7,628
A
registry
A+B
+ EHR
A+B+C
+ assays
+D
+ research
b
compute-optimal size (M parameters)
10^2
10^1
10^0
10^{-1}
ALLO-FM v1: 30–200 M
A+B
A
A+B+C lies on the A+B line:
the deep assays add information,
not tokens
10^3
10^4
10^5
patients in the pretraining corpus
c
A100-hours (log scale)
10^8
10^7
10^6
10^5
10^4
10^3
10^2
10^1
10^0
3 h
$6
115 h
$230
adapters, scaling,
ablations, 5 seeds
233,333 h
$0.5M
one run
A+B, 100k pts
whole
programme
7B LLM
(for scale)