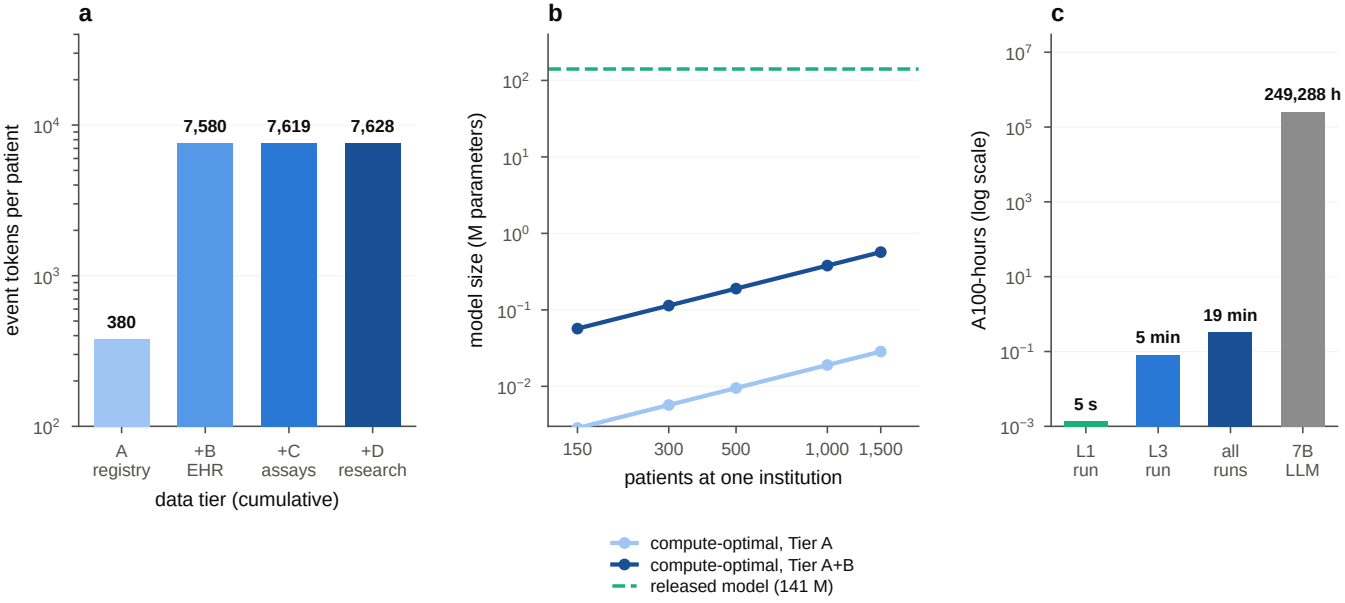
a
event tokens per patient
10^4
10^3
10^2
380
7,580
7,619
7,628
A registry
+B EHR
+C assays
+D research
data tier (cumulative)
b
model size (M parameters)
10^2
10^1
10^0
10^{-1}
10^{-2}
150
300
500
1,000
1,500
patients at one institution
compute-optimal, Tier A
compute-optimal, Tier A+B
released model (141 M)
c
A100-hours (log scale)
10^7
10^5
10^3
10^1
10^{-1}
10^{-3}
5 s
5 min
19 min
249,288 h
L1 run
L3 run
all runs
7B LLM